Wenjie Luo

AI Product Manager | **Graduating Jan 2027**

Singapore | +65 8084 8140 | +86 183 1949 6086 | e1576499@u.nus.edu
github.com/LUOaini1213 | luoaini1213.github.io (demos, source and evaluation reports)

PROFILE

- **Reliable agents:** the LLM routes, tools compute numbers, people approve risky actions, evaluations decide.
- **Metrics first:** acceptance criteria up front; features kept or cut by ablations and shadow evaluation.
- **Cross-market:** overseas VOC at Chery, a 7-market review study, Singapore site work; English and Chinese.

EDUCATION

National University of Singapore | MSc Civil Engineering (Transportation) Aug 2025 - Jan 2027

Coursework: ML Foundations (CEE5901QA), Advanced AI for CEE (CEE5901QB), intelligent transport and simulation.

Chang'an University | BEng Transportation Engineering Sep 2021 - Jun 2025

GPA 3.16/5.0, approximately top 15% of the cohort; class Youth League branch secretary.

EXPERIENCE

Far East Facade | Project Engineer, Engineering Department Jun 2026 - Present

- **Materials and suppliers:** on a semiconductor-fab façade project, track 15 material types on a status board and coordinate sampling and third-party testing with 5 suppliers.
- **Tooling:** use Python to extract delivery and test fields from supplier emails, replacing a hand-copied ledger.

Chery International | Product Intern May 2026 - Jun 2026

- **Customer insight:** ran competitor win/loss and complaint analysis; built **15** personas and **5** VOC reports for product definition and facelifts.
- **Market and schedule:** tracked target-market sales with bubble charts and KD localisation milestones on a Gantt chart.

Caixun Technology | Product Intern Jun 2025 - Aug 2025

- Independently wrote 3 PRDs and prototyped 5 core pages in Axure, all scheduled for development; moved a hard-to-find prize-draw entry to a side pop-up and followed retesting, resolving 10+ issues.

PROJECTS

Civil Buddy | Engineering Agent Workspace (open source) HITL Evaluation MCP Jul 2026 - Present

- **Scope:** 16 categories, 66 job roles; only deterministic tools produce numbers and risky actions need human confirmation; adds a Rust workbench with an MCP entry and a CAD-to-3D preview.
- **Packing loop:** material input → tool-computed container counts and coordinates → human confirmation → 3D load plan; PRD and acceptance checklist; all **128** runs completed, 57 judged not to fit and returned to a person.

Counterask | Conversational WebMCP Storefront Clarifying questions Checkout approval Sep 2026

- **Interaction:** when an answer would be a guess, search asks one clarifying question; the agent can fill checkout, but only the user can place the order (verified in Chrome 152).
- **Evaluation:** 9,901 real products, 800 queries; Hit@10 rose from 0.793 to **0.999** (0.991 on held-out phrasings); refusals read as requirements fell from 100% to 0 and budget overruns from 31% to 0.

Keeta Review Analysis | Cross-market Customer Feedback (public data) Sep 2026

- Analysed 2,572 App Store reviews across 7 markets; the dashboard shows low-star themes by market. Compared multilingual rules with local-LLM labels; SQL metrics are recomputed in pandas and checked in CI.

Malaysia Automotive Data Assistant | Metric-first Q&A Aug 2026

- Separates wholesale (TIV) from registration figures and stops to ask when the metric is unclear; on 30 authored questions it answered 22 correctly and asked for clarification on 8.

SKILLS

Product	requirements, PRDs and acceptance criteria, Axure, VOC analysis, personas, competitor analysis
AI products	agent tool calling, human-in-the-loop approval, RAG, retrieval evaluation, ablations, observability
Data and tools	Python, SQL, pandas, TypeScript, SUMO. Languages: Chinese (native), English (working)