Wenjie Luo

AI Product Manager | Graduating Jan 2027

Singapore | +65 8084 8140 | +86 183 1949 6086 | e1576499@u.nus.edu
github.com/LUOaini1213 | luoaini1213.github.io (demos and evaluation reports)

PROFILE

- Build reliable agentic AI products: the LLM orchestrates, deterministic tools compute the numbers, people approve high-risk actions, and reproducible evaluations back every claim.
- Experience across PRDs and prototypes, overseas-market VOC and competitor research, and multi-party coordination on Singapore construction projects; working in English and Chinese.

EDUCATION

National University of Singapore | MSc Civil Engineering — Aug 2025 - Jan 2027
Transport focus; coursework: ML Foundations, Advanced AI for CEE, intelligent transport and simulation.

Chang'an University | BEng Transportation Engineering — Sep 2021 - Jun 2025
GPA 3.16/5.0; approximately top 15% of the cohort.

EXPERIENCE

Far East Facade | Project Engineer, Engineering Department — Jun 2026 - Present
- Semiconductor-fab façade project: maintain the list and delivery plan for 15 material types on a status board; coordinate sampling and third-party testing with 5 suppliers.
- Use Python to extract delivery and test fields from supplier shipment and test emails, replacing a hand-copied ledger.

Chery International | Product Intern — May 2026 - Jun 2026
- Ran competitor win/loss and complaint analysis; built 15 personas and 5 VOC reports for product definition and facelifts; tracked overseas target-market sales with length, price and registration bubble charts.

Caixun Technology | Product Intern — Jun 2025 - Aug 2025
- Independently wrote 3 PRDs and prototyped 5 core pages in Axure; proposals were accepted and scheduled. Moved a hard-to-find prize-draw entry to a side pop-up and followed development and retesting, resolving 10+ issues.

Expeditors | Supply Chain Assistant — Dec 2024 - Apr 2025
- Worked with carriers, clients and customs brokers in English; streamlined Excel document checks, improving efficiency by about 30%.

PROJECTS

Civil Buddy | Engineering Agent Workspace (open-source flagship) — Jul 2026 - Present
- Workspace with 66 role skills; packing flow: material input → tool-computed container counts and coordinates → human confirmation → 3D load plan. The model never writes the numbers.
- Wrote the packing PRD and acceptance checklist (refusal reasons, recovery, cost limits); all 128 evaluation runs completed, 57 judged not to fit and handed back to a person.

Counterask | Conversational WebMCP Storefront — Sep 2026
- Search asks one clarifying question when an answer would be a guess; checkout is a declarative form the agent can fill, but only the user can place the order (verified in Chrome WebMCP).
- 9,901 real products, 800 queries: Hit@10 rose from 0.793 (keyword matching) to 0.999 (0.991 on held-out phrasings); negation errors fell from 100% to 0 and budget overruns from 31% to 0.

Glass Box | Agent Observability Middleware (TechJam 2026) — Sep 2026
- Turns each agent run into an OpenTelemetry GenAI span waterfall that pinpoints failing steps; shows OTLP delivery status, passes 503, timeout and killed-process recovery tests, and redacts secrets before persistence.

Keeta Review Analysis | Cross-market Customer Feedback — 2026
- Analysed 2,572 public App Store reviews across 7 markets; dashboard of low-star themes by market. Compared multilingual rules with local-LLM labels; SQL metrics recomputed in pandas and checked in CI.

Malaysia Automotive Data Assistant — Aug 2026
- Separates wholesale (TIV) from registration figures and stops to ask when the metric is ambiguous or data is missing; on 30 authored questions it answered 22 correctly and asked for clarification on 8.

More | Evaluation and Inference Performance — 2026
- LLM-judge study auditing reference-answer, position and phrasing effects; traced a vLLM throughput collapse on a GTX 1650 to a cuBLAS fp16 GEMM small-M cliff (99.5 → 4.3 GiB/s; absent on T4) and filed issues with PyTorch and vLLM.

SKILLS

Product: requirements, PRDs and acceptance criteria, Axure, VOC analysis, personas, competitor analysis
AI products: agent tool calling, human-in-the-loop approval, RAG, evaluation and ablations, OpenTelemetry
Data and tools: Python, SQL, pandas, TypeScript, SUMO. **Languages:** Chinese (native), English (working)